



Blockchain Principles and Applications

Amir Mahdi Sadeghzadeh, PhD

Data and Network Security Lab (DNSL)
Trustworthy and Secure AI Lab (TSAIL)

Basic Cryptographic Primitives

1. Cryptographic Hash Functions

2. Hash Accumulators
Merkle trees

Centralized Blockchain

3. Digital Signatures


Decentralized Blockchain

Cryptography Background

(1) cryptographic hash functions

An efficiently computable function $H: M \rightarrow T$

where $|M| \gg |T|$



Collision resistance

Def: a collision for $H: M \rightarrow T$ is pair $x \neq y \in M$ s.t. $H(x) = H(y)$

$|M| \gg |T|$ implies that many collisions exist

Def: a function $H: M \rightarrow T$ is collision resistant if it is “hard” to find even a single collision for H (we say H is a CRF)


Example: **SHA256:** $\{x : \text{len}(x) < 2^{64} \text{ bytes}\} \rightarrow \{0,1\}^{256}$

(output is 32 bytes)

Hash Functions

Defining Properties:

1. Arbitrary sized inputs
2. Fixed size deterministic output
3. Efficiently computable
4. Minimize collisions



Canonical application:

Hash Tables

Store and retrieve data records

Application: committing to data on a blockchain

Alice has a large file m . She posts $h = H(m)$ (32 bytes)

Bob reads h . Later he learns m' s.t. $H(m') = h$

H is a CRF \Rightarrow Bob is convinced that $m' = m$
(otherwise, m and m' are a collision for H)

We say that $h = H(m)$ is a **binding commitment** to m

(note: not hiding, h may leak information about m)

Committing to a list (of transactions)

Alice has $S = (m_1, m_2, \dots, m_n)$

32 bytes



Goal:

- Alice posts a short binding commitment to S , $h = \text{commit}(S)$
- Bob reads h . Given $(m_i, \text{proof } \pi_i)$ can check that $S[i] = m_i$
Bob runs $\text{verify}(h, i, m_i, \pi_i) \rightarrow \text{accept/reject}$

security: adv. cannot find (S, i, m, π) s.t. $m \neq S[i]$ and
 $\text{verify}(h, i, m, \pi) = \text{accept}$ where $h = \text{commit}(S)$

Cryptographic Hash Function

Extra Property:

Specialized one way function

Canonical application:

Puzzle generation
mining process

**Hash(nonce, block-hash) <
Threshold**

Another application: proof of work

Goal: computational problem that

- takes time $\Omega(D)$ to solve, but
- solution takes time $O(1)$ to verify

(D is called the **difficulty**)

How? $H: X \times Y \rightarrow \{0, 1, 2, \dots, 2^n - 1\}$ e.g. $n = 256$

- puzzle: input $x \in X$, output $y \in Y$ s.t. $H(x, y) < 2^n/D$
- verify(x, y): accept if $H(x, y) < 2^n/D$

Another application: proof of work

Thm: if H is a “random function” then the best algorithm requires D evaluations of H in expectation.

Note: this is a parallel algorithm

⇒ the more machines I have, the faster I solve the puzzle.

Proof of work is used in some consensus protocols (e.g., Bitcoin)

Bitcoin uses $H(x, y) = \text{SHA256}(\text{SHA256}(x.y))$

Hash Pointer

Hash of the information acts as pointer to location of information

Regular pointer: retrieve information

Hash pointer: retrieve information and verify the information has not changed

Regular pointers can be used to build data structures: linked lists, binary trees.

Hash pointers can also be used to build related data structures. Crucially useful for blockchains. In fact, **blockchain itself is a hash pointer-based data structure.**

Blockchain: a linked list via hash pointers

Block: Header + Data

Header: Pointer to previous block
= hash of the previous block

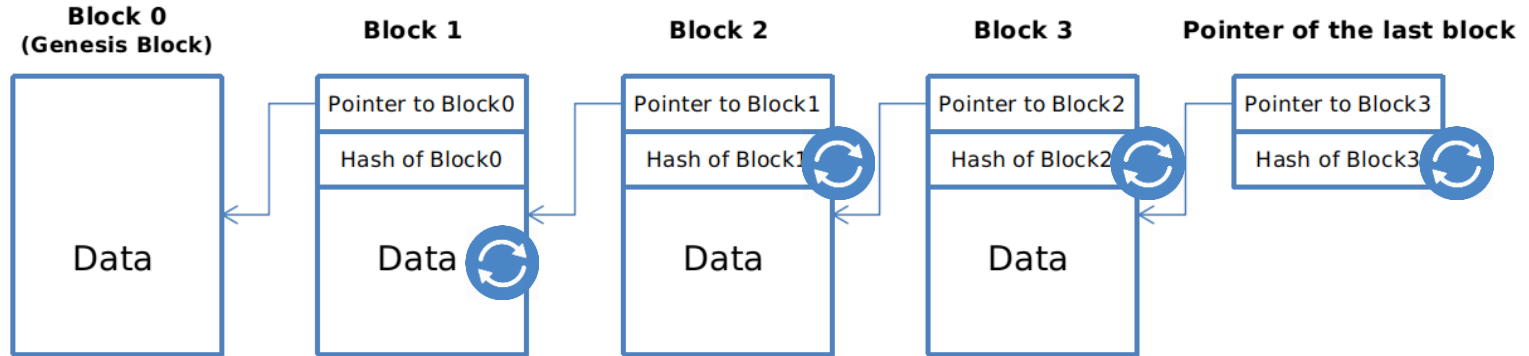
Data: information specific to the block

Application: tamper evident information log

Head of the chain being known is enough to find tamper evidence in any internal block

Hence the phrase: **block chain**
Or simply: **blockchain**

Blockchain: a linked list via hash pointers



Allows the creation of a tamper-evident information log

How about searching for specific data elements?

Merkle Tree

Binary tree of hash pointers

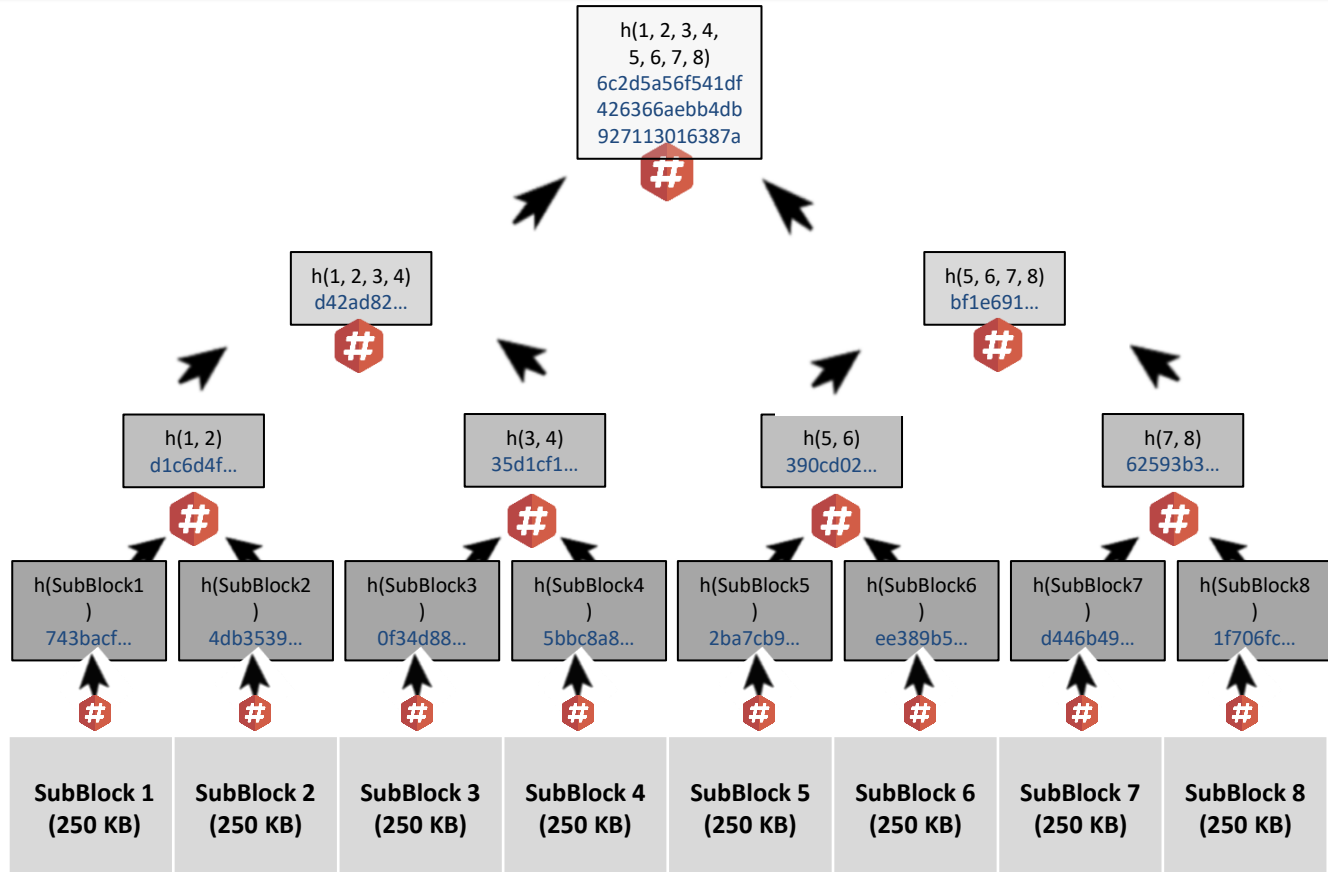
Retain only the root of the tree

Tamper of any data in the
bottom of the tree is evident

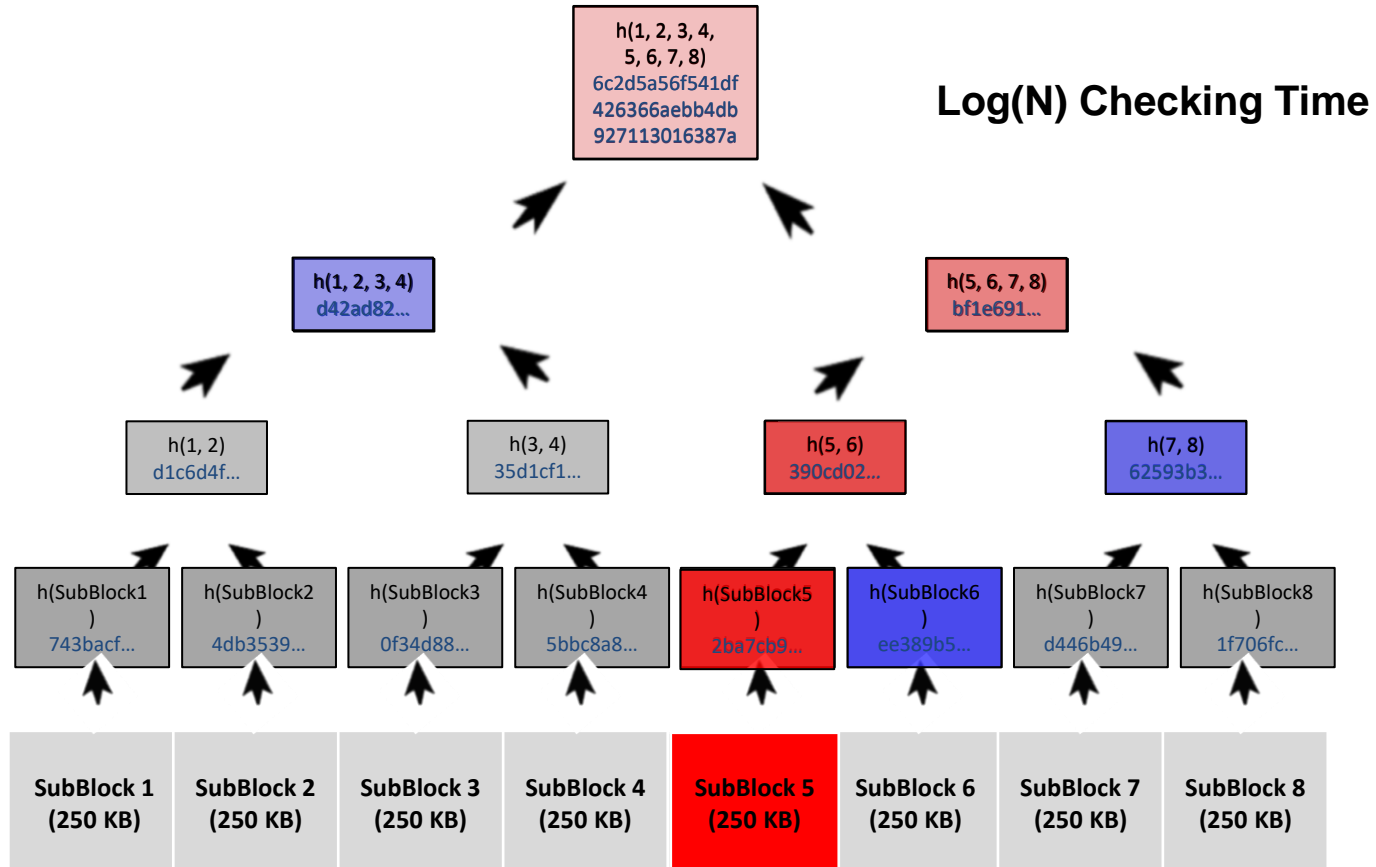
Proof of Membership

Proof of Non-membership

Merkle Tree



Proof of Membership



Merkle tree

(Merkle 1989)

commitment

h

Merkle tree
commitment

m_1 m_2 m_3 m_4 m_5 m_6 m_7 m_8

list of values S

Goal:


- commit to list S of size n
- Later prove $S[i] = m_i$

Merkle tree

(Merkle 1989)

[simplified]

commitment



Goal:

- commit to list S of size n
- Later prove $S[i] = m_i$

To prove $S[4] = m_4$,
proof $\pi = (m_3, y_1, y_6)$

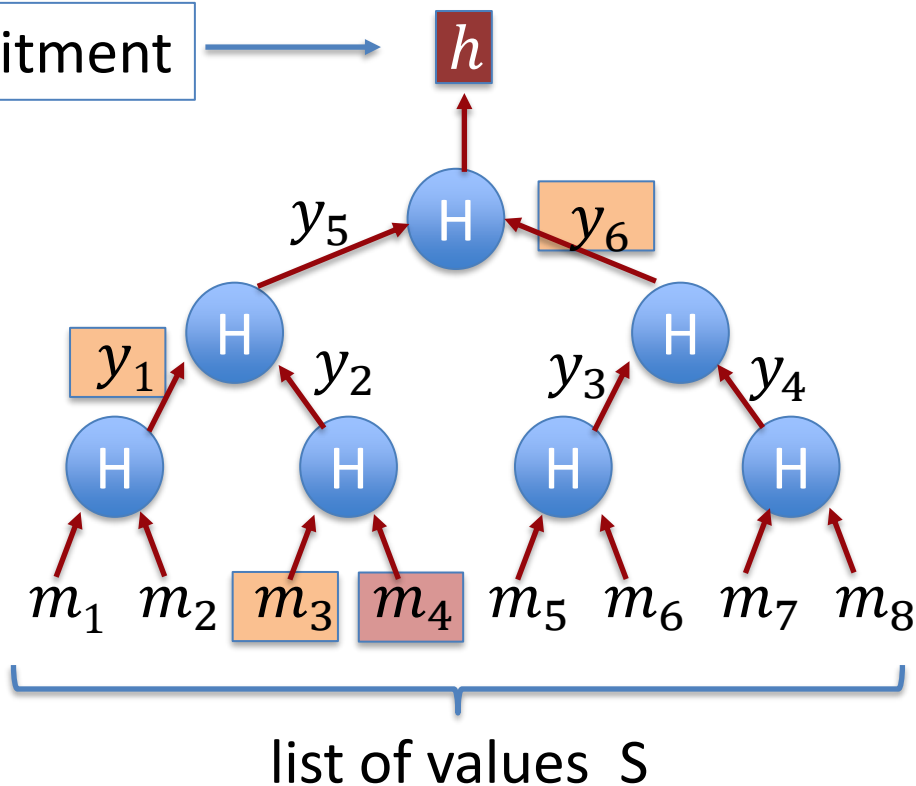
length of proof: $\log_2 n$

Merkle tree

(Merkle 1989)

[simplified]

commitment



To prove $S[4] = m_4$,
proof $\pi = (m_3, y_1, y_6)$

Bob does:

$$y_2 \leftarrow H(m_3, m_4)$$

$$y_5 \leftarrow H(y_1, y_2)$$

$$h' \leftarrow H(y_5, y_6)$$

accept if $h = h'$

Merkle tree

(Merkle 1989)

Thm: For a given n : if H is a CRF then

adv. cannot find (S, i, m, π) s.t. $|S| = n$, $m \neq S[i]$,

$h = \text{commit}(S)$, and $\text{verify}(h, i, m, \pi) = \text{accept}$

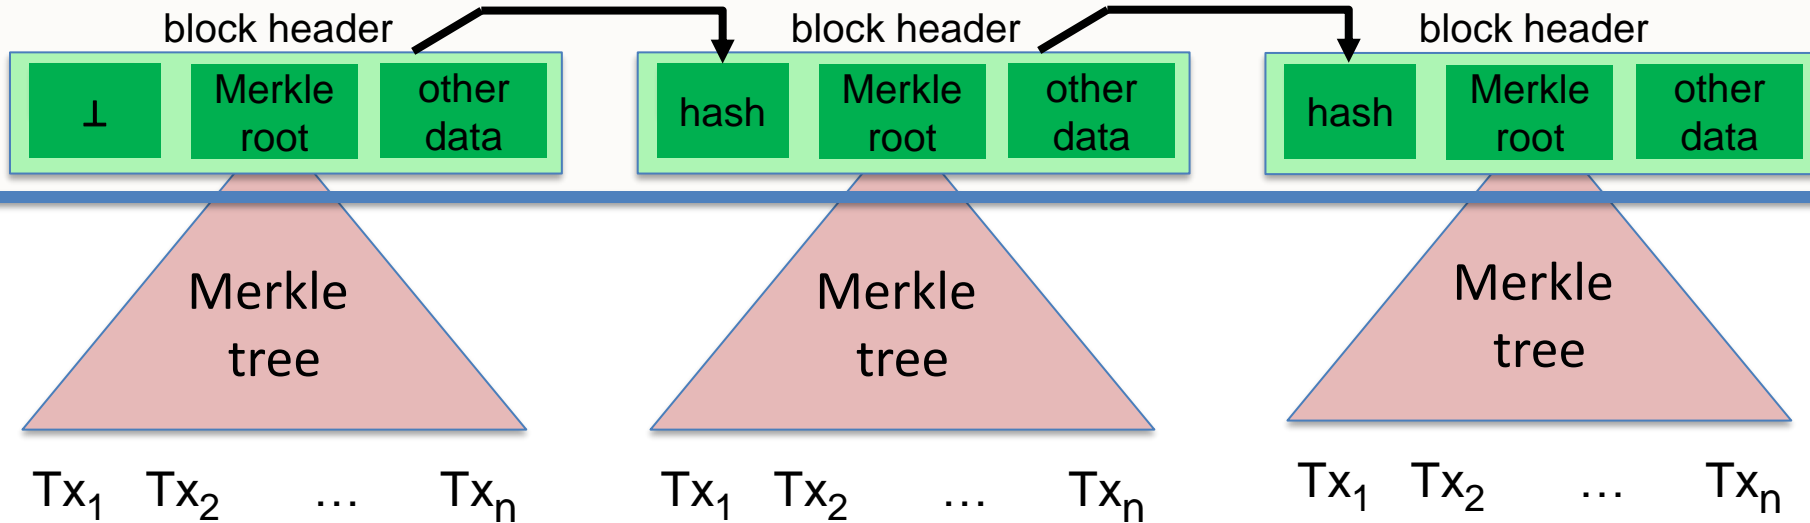
(to prove, prove the contra-positive)

How is this useful? To post a block of transactions S on chain suffices to only write $\text{commit}(S)$ to chain. Keeps chain small.

\Rightarrow Later, can prove contents of every Tx.

Abstract block chain

blockchain



Merkle proofs are used to prove that a Tx is “on the block chain”

Blockchain with Merkle Trees

Block: Header + Data


Header: Pointer to previous block
= hash of the previous block header and Merkle root of data of previous block

Data: information specific to the block

Application: Centralized tamper evident information log with efficient proof of membership of any data entry


Head of the chain being known is enough to find tamper evidence in any internal block

Blockchain with Merkle Trees



Signatures

Physical signatures: bind transaction to author

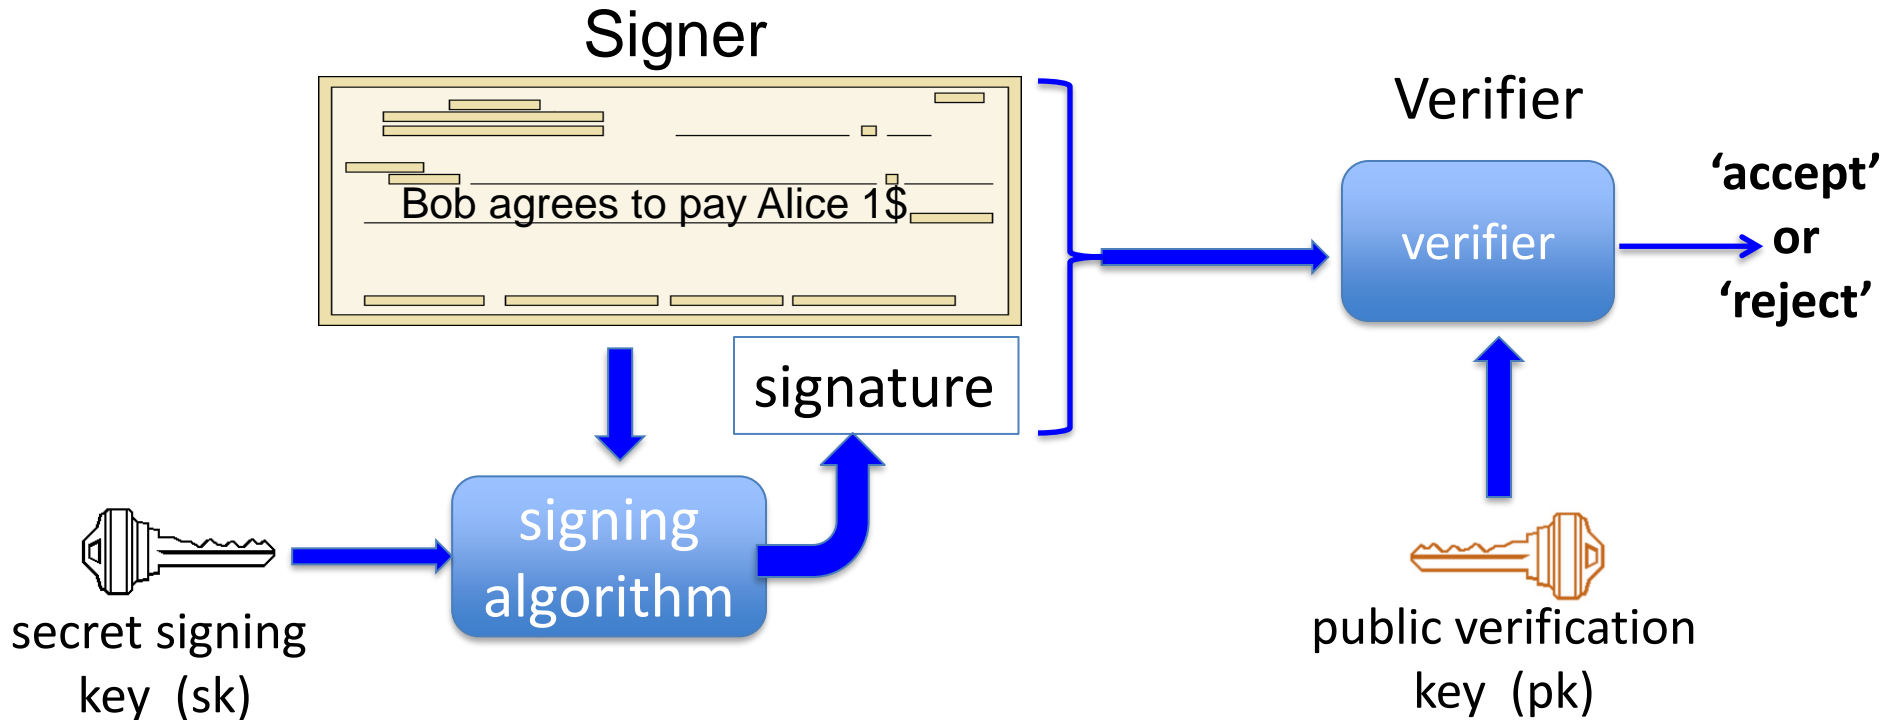


Problem in the digital world:

anyone can copy Bob's signature from one doc to another

Digital signatures

Solution: make signature depend on document



Digital signatures: syntax

Def: a signature scheme is a triple of algorithms:

- **Gen()**: outputs a key pair (pk, sk)
- **Sign**(sk, msg) outputs sig. σ
- **Verify**(pk, msg, σ) outputs 'accept' or 'reject'

Secure signatures: (informal)

Adversary who sees signatures **on many messages** of his choice, cannot forge a signature on a new message.


Families of signature schemes

1. RSA signatures (old ... not used in blockchains):
 - long sigs and public keys (≥ 256 bytes), fast to verify
2. Discrete-log signatures: Schnorr and ECDSA (Bitcoin, Ethereum)
 - short sigs (48 or 64 bytes) and public key (32 bytes)
3. BLS signatures: 48 bytes, aggregatable, easy threshold
(Ethereum 2.0, Chia, Dfinity)
4. Post-quantum signatures: long (≥ 600 bytes)

Signatures on the blockchain

Signatures are used everywhere:

- ensure Tx authorization,
- governance votes,
- consensus protocol votes.



Decentralizing the Blockchain

Digital Signatures

Decentralized Identity Management

Elements of a cryptocurrency

Digital Signatures

Key generation

(secretkey, publickey) =
Generatekeys(keysize)

Randomized function

Signature

Sig = *sign*(secretkey, message)

Verification

verify(publickey, Sig, message)

Unforgeable Signatures

Unforgeable

Computationally hard to generate a verifiable signature without knowing the secret key

ECDSA

Elliptic Curve Digital Signature Algorithms

Cryptographically secure against an adaptive adversary

Signatures in Practice

Elliptic Curve Digital Signature Algorithm
(ECDSA)

Standard part of crypto libraries

Public key: 512 bits

Secret key: 256 bits

Message: 256 bits

Note: can sign hash of message

Signature: 512 bits

Decentralized Identity Management

Public keys are your identity
address in Bitcoin terminology

Can create multiple identities
(publickey, secretkey) pairs
publish publickey
sign using secretkey

Can create oneself
verifiable by others

Resources

- ECE/COS 470, Pramod Viswanath, Princeton 2024
- CS251, Dan Boneh, Stanford 2023