



# Blockchain Principles and Applications

Amir Mahdi Sadeghzadeh, PhD

Data and Network Security Lab (DNSL)  
Trustworthy and Secure AI Lab (TSAIL)

# Recap

# Bitcoin latency

Time from when a transaction was broadcast until the transaction is confirmed in the ledger

- $\tau_1$ : Time from when a transaction was broadcast until the transaction is put into a mined block B
- $\tau_2$ : Time from when the transaction was put into a mined block B until block B is  $k$ -deep in the longest chain

$$\tau = \tau_1 + \tau_2$$

$\tau_2$  is the real bottleneck, depends on how large  $k$  is.

# Bitcoin latency

Assume low forking ( $\lambda\Delta \ll 1$ ),

$$\tau = \frac{k}{(1 - \beta)\lambda}$$

Depth of  
blocks

From Lecture 6, error probability

$$\epsilon = e^{-ck}$$

Block  
arrival rate

$$\tau = \frac{\frac{1}{c} \log(\frac{1}{\epsilon})}{(1 - \beta)\lambda} = O\left(\frac{1}{\lambda} \log\left(\frac{1}{\epsilon}\right)\right)$$

**Latency and security are coupled**

# Improve Bitcoin latency

Only way to improve latency is to

- reduce  $k$ ; but this reduces security
- Increase  $\lambda$ ; but this also reduces security

Ethereum:  $\frac{1}{\lambda} = 15\text{s}$ ;  $k = 100$

- latency = 25 minutes
- Way better than Bitcoin performance; improvement simply by picking better parameters.

# Improve Bitcoin latency

Question: can we make relatively small changes to the longest chain protocol and PoW mining while scaling latency?

Key Requirement:

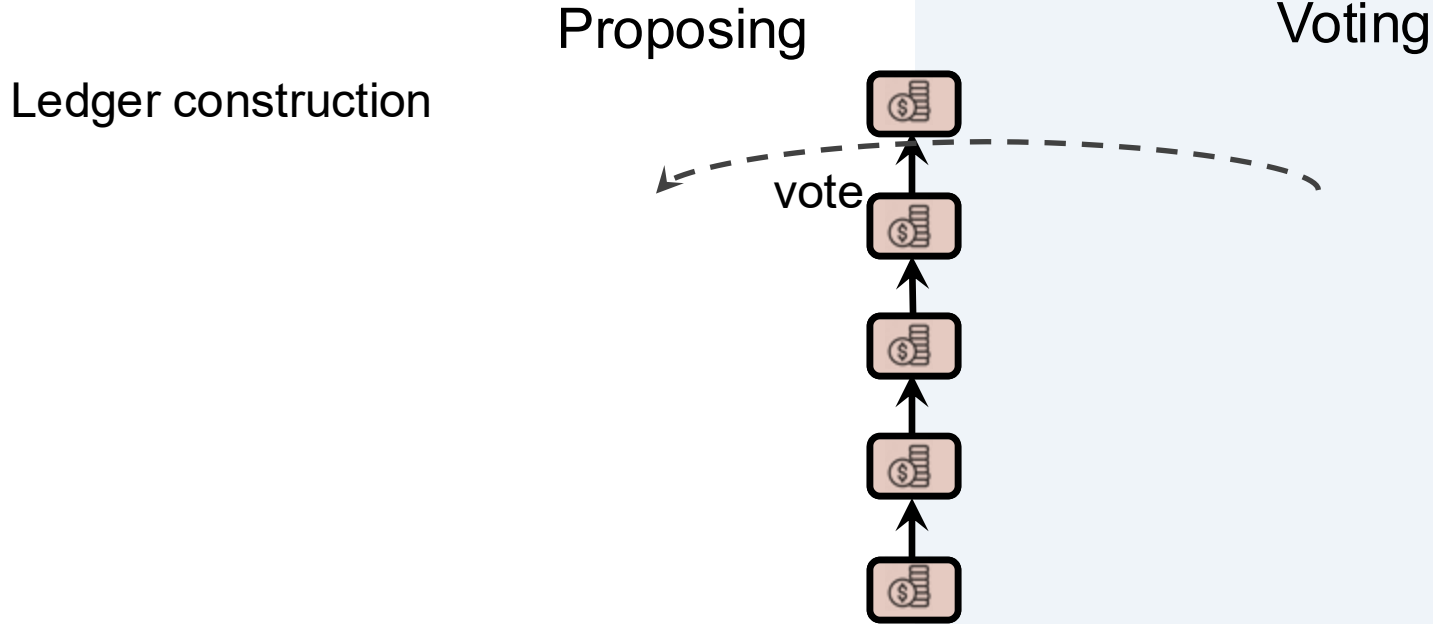
- Do not want latency to depend on security level
- **Decouple security from latency**

# Prism

Prism achieves optimal latency

- **Decoupling principle**: separate performance from security
- Prism 1.0 achieves optimal throughput; last lecture

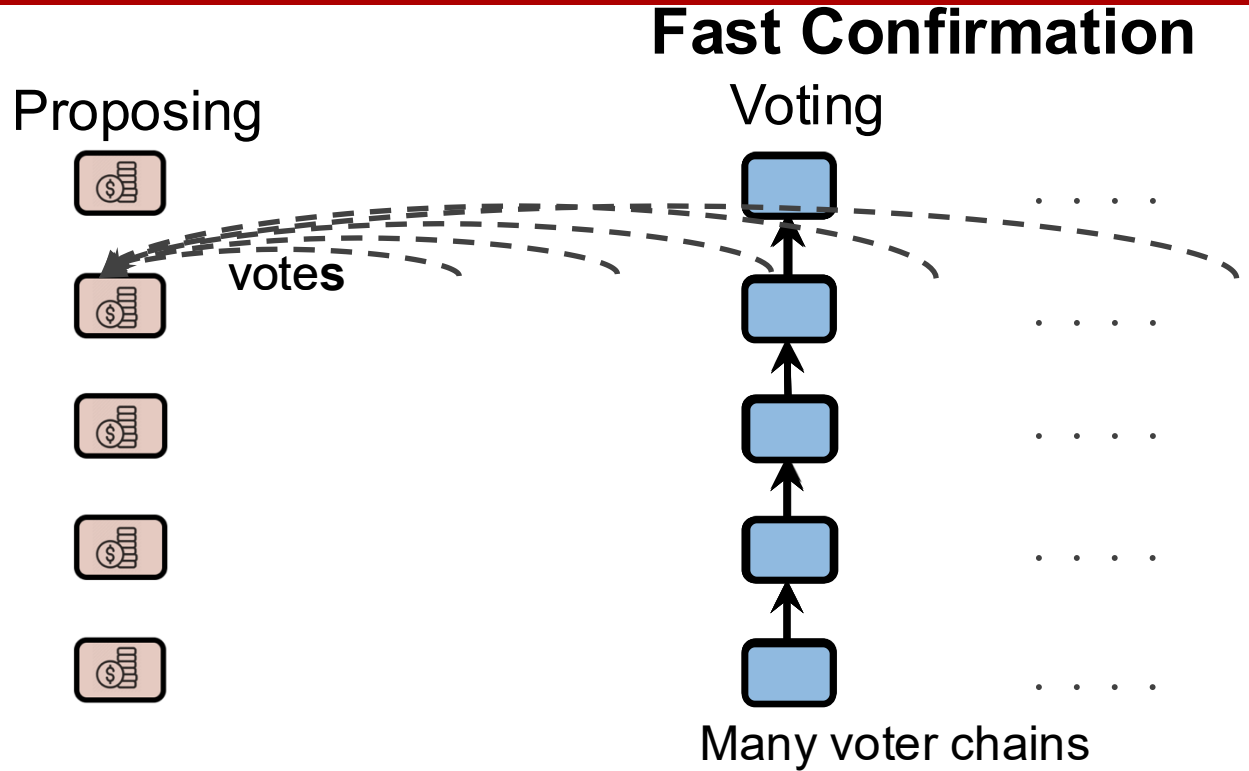
# Bitcoin → Deconstruct



1. Select **votes** along **longest** voter chain
2. Order the proposer blocks by votes




# Bitcoin → Deconstruct → Prism



Ledger Construction: For each level choose the proposer block with **maximum** votes

# Sharding

# Blockchains & Full replication



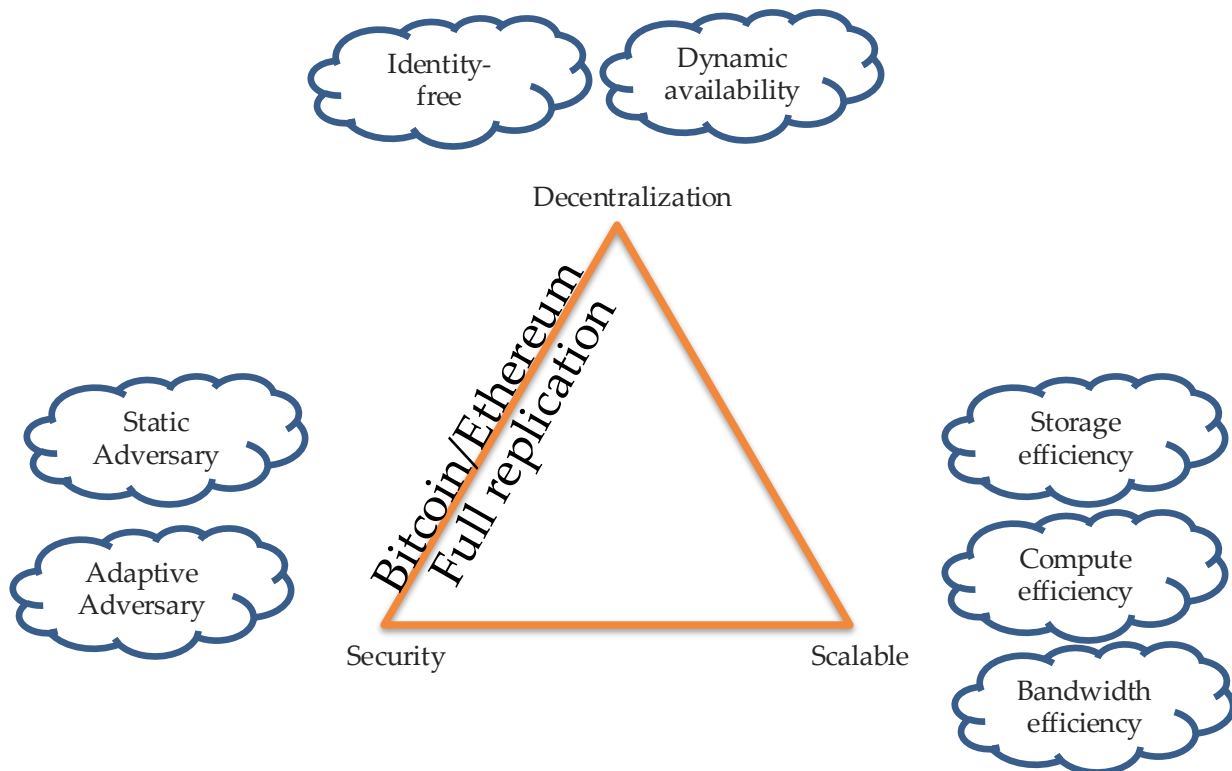
# Full replication – Consensus problems

- Throughput of consensus decreases with an increase in size of the number of participating nodes
- For Nakamoto delay  $\propto O(\log(N))$
- For Nakamoto, communication load =  $O(N^2)$

# Full replication – resource usage


- All nodes process the same transactions
- **Communication:** The transaction has to traverse the complete network at least once
- **Storage:** All nodes have to store the complete state, the account details of everyone!
- **Compute:** All nodes have to validate all transactions and update the ledger every block

# Trilemma




# First Approach– Maintain multiple blockchains

- Divide ledger into K shards
- Each shard is a separate blockchain



Complete ledger



K Shards


# Problems

1. Reduced security – An adversary can concentrate on one shard
2. How to transfer money from one shard to another?
3. If such transfer is possible, adversary can transfer non-existent funds from infected shard to non-infected shard.




# Node to shard allocation (N2S)

- Extension of the first order approach
- Allocates each consensus nodes to one shard randomly




# Multiconsensus

- Each shard runs its own consensus



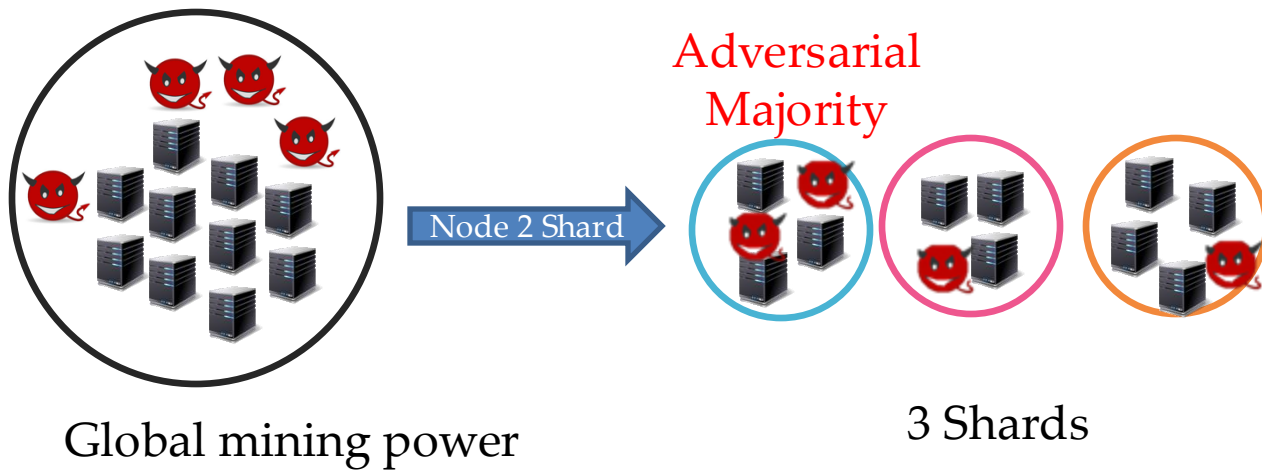
# Scaling: $O(K)$

- Nodes only maintain the state of their own shard



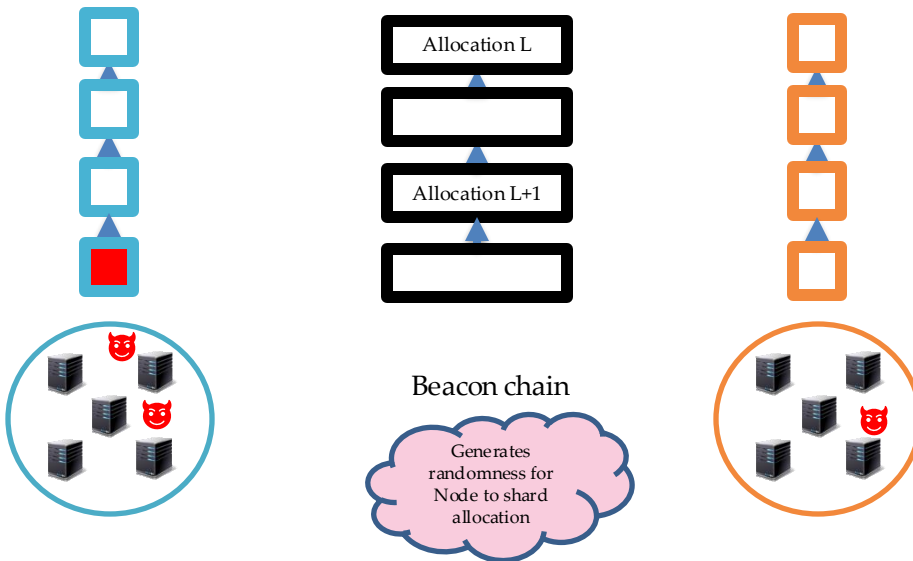
# Drawback 1: Proportional representation

- Need large number of nodes per shard to ensure honest majority in a shard



# Drawback 2: Security

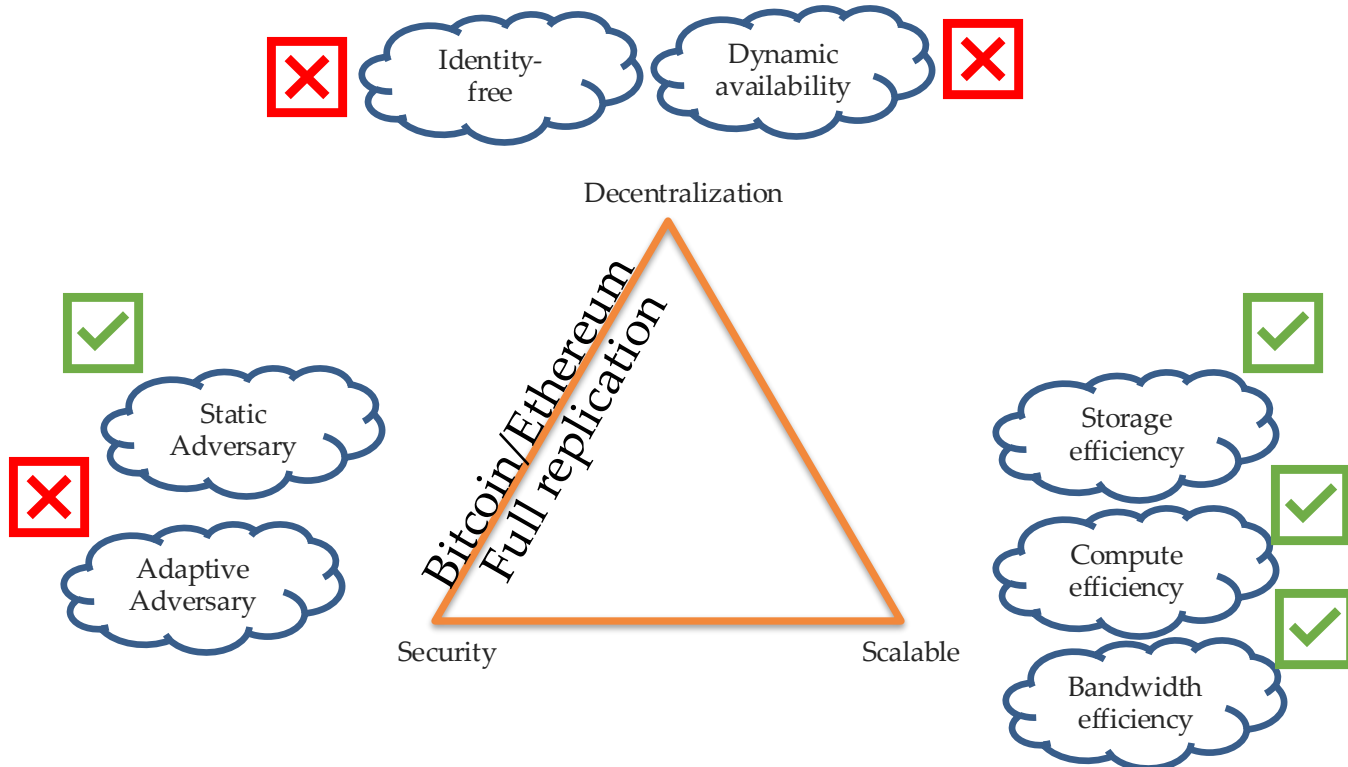
- Not resilient to  $O(1/K)$  adaptive adversary



# Drawback 3: Node identity

- Existing works vary in the implementation of Node to shard allocation
  - Different ways of randomness generation
  - Different rate of re-allocation
- All Node to shard allocation algorithms require consensus node identity

# Trilemma revisited




# Identity-free sharding

- If the protocol is identity free, we cannot use Node to Shard allocation algorithm
- Only choice is to allow nodes to **self-allocate**
- Adversaries can congregate on one shard; safety and liveness can be easily broken
- Solution: **Uniconsensus** architecture



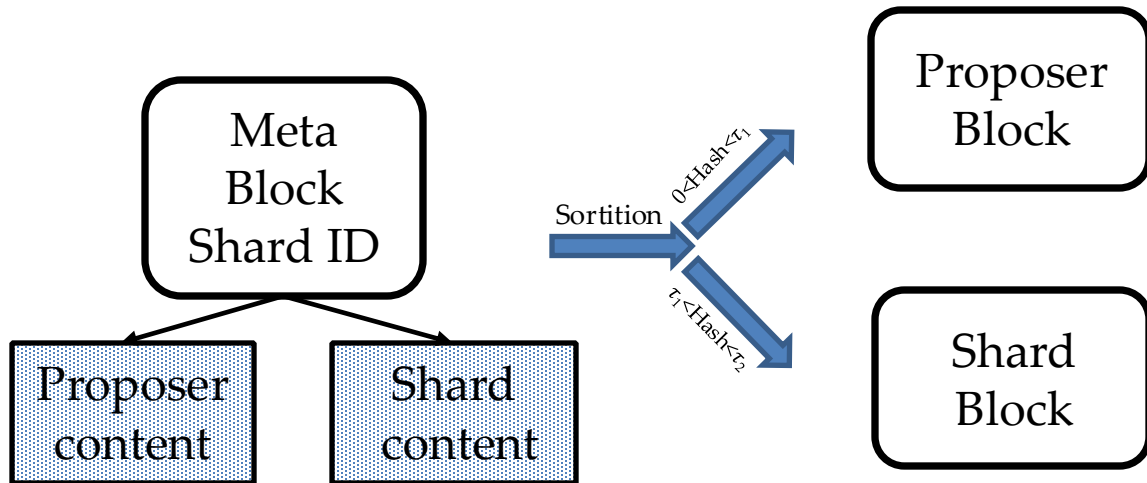
# Uniconsensus: Decoupled validation

- Shard transaction ordering can be decoupled from validation
- Extension of the deconstruction ideas from Fruitchains and Prism

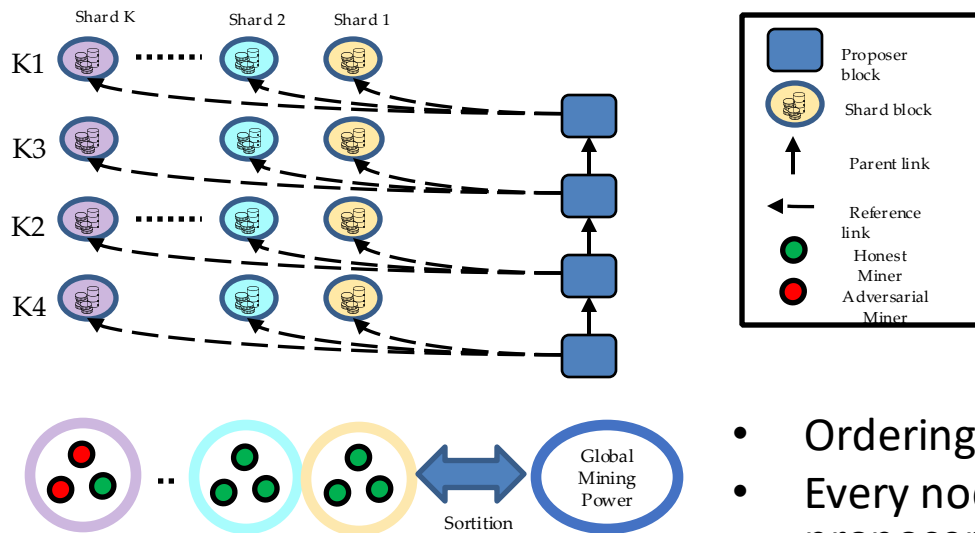


# Uniconsensus: Sortition

- A node can maintain any shard of its choice
- It will maintain an ordering chain in parallel
- All nodes mine shard block and proposer block together



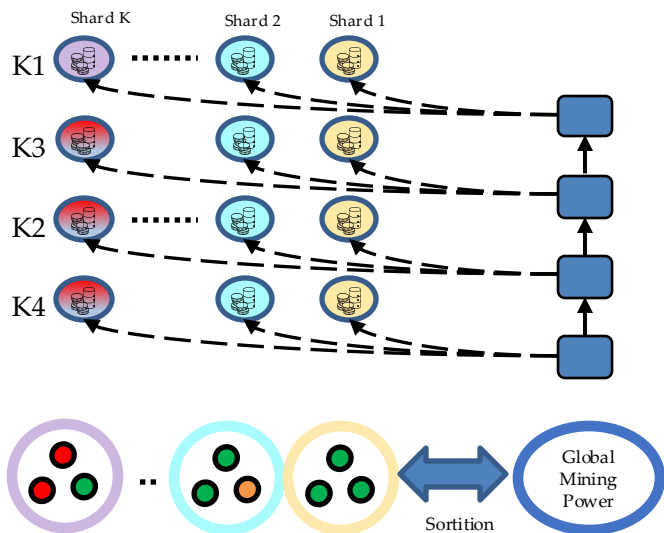
# Uniconsensus: Safety



- Ordering: **Proposer chain**
- Every node maintains proposer chain
- Adversarial majority in a shard does not violate safety


# Uniconsensus: Liveness

Chain-quality visualization



- Adversaries congregate: drown out honest miners
- Dynamic self allocation can prevent such attacks

# Uniconsensus Architecture




- Safety shared
- Liveness is sacrificed

# Dynamic Self-allocation



- Honest nodes **adapt** to adversaries and allocate themselves to new shards
- We want honest nodes to (re)allocate themselves to shards under attack

# Simple DSA


- Adapt to the adversary by following the adversary's last move
- Assume that the adversary's past allocations are known



# Practical Implementation



- Honest miners estimate total mining power from proposer chain
- Honest mining power allocation can be estimated from DSA policy





# Resources

- ECE/COS 470, Pramod Viswanath, Princeton 2024